

SHRUTI JAIN

+1-412-589-5683 | Pittsburgh, PA | spjain@andrew.cmu.edu | [LinkedIn](#) | [Personal Website](#)

EDUCATION

Carnegie Mellon University | Master of Science in Computer Vision Aug 2025 - Dec 2026(Expected)
Coursework: Visual Learning and Recognition, Advanced Computer Vision, Introduction to Machine Learning
CGPA: 4.11/4

Manipal Institute of Technology | Bachelor of Technology in Information Technology 2019-2023
CGPA: 8.71/10

SKILLS

Programming Languages/Tools: Python, C++, Java, Docker, Kubernetes, SQL
ML and CV Frameworks/Libraries: PyTorch, OpenCV, ROS, CUDA, scikit-learn, TensorFlow, Matplotlib
AI Topics: Computer Vision, Transformers, VLMs, AI Agents, Generative Models, Object Detection, Image Segmentation, Zero-shot Learning, Semi-supervised Learning, Domain Adaptation, AI evaluation

RELEVANT EXPERIENCE

FlowState AI | *Computer Vision Research Intern* May 2026 - Aug 2026

- [\[Slides\]](#) **Architected and implemented** an end-to-end **AI video-editing system** that converts long-form footage into prompt-driven, story-based, intelligently re-framed (e.g., from 16:9 to 9:16) content.
- Developed shot-aware subject grounding, tracking, and stable framing logic to identify salient subjects and preserve cinematographic composition with workflow orchestration and self-evaluation.

AIRLab - Carnegie Mellon University | *Research Assistant* Jan 2026 - April 2026

- [\[Project Website\]](#) Developing a multimodal medical triage pipeline for the **DARPA Triage Challenge** using RGB and thermal UAV imagery to predict 50+ hierarchical medical labels, under the guidance of **Prof. Sebastian Scherer**.
- Built a two-stage system combining conditional diffusion-based RGB-to-IR generation and VLM fine-tuning with LoRA/GRPO optimization to improve triage reasoning under limited paired RGB-IR data.

Carnegie Mellon University, USA | *Research Assistant* Aug 2025 – Dec 2025

- Worked under **Prof. Deva Ramanan** on developing DetPO, a gradient-free black-box prompt optimization framework achieving up to 9.7% improvement over prior approaches on Roboflow20-VL and LVIS.

Palo Alto Networks | *Software Developer* Oct 2023 - April 2025

- Worked on features and APIs of the microservices (**Java Spring Boot**) for ingesting events from tenants' cloud accounts (**AWS, Azure, GCP**) for near real-time threat detection using **Artificial Intelligence**.
- Integrated a **scheduler-worker service** that schedules jobs to poll cloud account for events with a **job management framework** based on **Amazon SQS** and **Redis cache**.

Robotics Research Centre (RRC) - IIIT Hyderabad | *Research Assistant* June 2023 – Sept 2023

- Achieved **14.4%** average decrease in the mean absolute scores across 3 datasets as compared to the SOTA for **Image Segmentation in self-driving cars** in adverse weather conditions.
- Implemented a **differentiable visual and latent prompting mechanism** to fine-tune on **Segment Anything Model** via **adaptors** through learnable image pre-processing blocks.

SELECTED PUBLICATIONS

- DetPO: In-Context Learning with Multi-Modal LLMs for Few-Shot Object Detection [\[Paper\]](#) [\[Website\]](#) [\[Code\]](#)
Published in ECCV 2026
- DiffPrompter: Differentiable Implicit Visual Prompts for Semantic-Segmentation in Adverse Conditions [\[Paper\]](#) [\[Project Page\]](#) [\[Code\]](#) **Published in IROS 2024**
- Light Weight Character and Shape Recognition for Autonomous Drones [\[Paper\]](#)
Workshop acceptance in ICML 2022
- A Studious Approach To Semi-Supervised Learning [\[Paper\]](#) [\[Code\]](#) **Workshop acceptance in NeurIPS 2021**

RELEVANT PROJECTS

- **Gated Multimodal Web Navigation Agents** - [\[Poster\]](#) A gated multimodal web navigation framework combining MindAct and Qwen2-VL on the Multimodal Mind2Web benchmark, improving robustness to adversarial popup attacks.
- **Editing 3D Gaussian Splats** - [\[Slides\]](#) Edits **3D Gaussian Splats** based on text prompts by using **Diffusion** based 2D Image Editing techniques.
- **Flash Attention for ViT** - [\[Poster\]](#) Implemented FlashAttention-1 and 2 in MiniTorch as custom CUDA autograd ops with tiled SRAM compute and online softmax.
- **Real-time Failed 3D Print Detector** - Detects failed 3D prints in real time using **OpenCV** and anomaly detection.
- **AUVSI SUAS Competition** [\[link\]](#) - Achieved Autonomous detection and classification of shapes (**55% accuracy**), letters (**75% accuracy**) and their colors from a rescue drone using **OpenCV** and **machine learning**.